
Small Language Models and the Synthetic Middle

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 This paper explores Small Language Models (SLMs) as reflexive tools that push
2 against dominant values embedded in large-scale language modeling systems.
3 Instead of optimizing for values like model objectivity, SLMs usefully distill
4 ideological constructions and social biases within language data. To illustrate
5 this effect, this project uses SLMs to examine the polarized discourse around
6 gender rights in the USA. Rather than smooth over conflicting perspectives about
7 gender, SLMs aggregate areas of ideological disagreement and misunderstanding
8 to reveal unexpected overlap between liberal and conservative views. By surfacing
9 intersections between politically opposed perspectives, SLMs make visible the
10 biases and structures of disagreement that large-scale models tend to obscure.

11 1 Toward a Synthetic Middle

12 Text generation tasks are often used for ideation, brainstorming, and drafting, as well as summarizing,
13 critiquing, and proofreading. These uses are centered on productive activities which harness the
14 generative affordances of machine learning (ML) to create new content.

15 This paper, by contrast, examines the *reflexive* potential of text generation—that is, how this mechanism
16 might surface latent structures, such as biased perspectives, from training data. It uses Small Language
17 Models (SLMs), derivations of GPT-2, to study the issue of gender rights in the USA from politically
18 polarized sources, what Bornschier et al. (2021) refers to as "far right and new left" views.

19 Due to their small size and under-developed status, SLMs do not merely reproduce viewpoints. Rather,
20 they tend to collapse opposing perspectives into apparently univocal declarations. Their outputs
21 lead to paper's key insight: that text generation produces aggregated language forms, compressing
22 multiple, often conflicting perspectives into a single statement—expressing what I call the *synthetic*
23 *middle*.

24 While a larger model might smooth over or resolve conflicting perspectives, the SLM's illusion
25 of a unified voice reveals traces of unexpected ideological overlap and intersection. These traces
26 of overlap and intersection contribute to academic debates in the fields of Feminist Studies and
27 Political Science, from theorizing about the widespread fears of "gender ideology" to recent trends in
28 electoral behavior and polarization (Butler, 2024; Bornschier et al., 2021). In doing so, this project
29 answers calls to operationalize ML as a tool for critique that both draws from and contributes to
30 interdisciplinary knowledge (Wang, 2025; Birhane et al., 2022; Devinney et al., 2022).

31 Crucially, this use of SLMs resists major drivers in current ML research and development. First,
32 it sidesteps the development goal of model objectivity which, as I explain below, overlooks how
33 language modeling embeds useful data for studying polarization in popular discourse. By working
34 with smaller, more controlled datasets, SLMs transform polarized views into sites of productive
35 intersection and potentially, intervention. Second, SLMs contribute to efforts for low-resource
36 computing, resisting the extractive "bigger is better mentality" which chases ever-increasing model

37 sizes (Varoquaux et al., 2025). This project shows how even underdeveloped, small models offer
38 powerful analytical possibilities.

39 2 On Objectivity

40 Birhane et al. (2022) identify values that drive ML development, including "Performance, Generaliza-
41 tion, Quantitative evidence, Efficiency, Building on past work, and Novelty". This project interrogates
42 another related value of ML research, the value of objectivity in model outputs.

43 LLM developers often frame objectivity as a driving goal for model development. OpenAI's "Model
44 Spec", which outlines the intended behavior for OpenAI language models, prescribe models that
45 "Assume an objective point of view," which includes "fairly describing significant views... and
46 allocat[ing] attention proportionally to their level of acceptance and evidential support" (OpenAI,
47 2025b).

48 Recently, OpenAI has made a concerted effort to reduce political bias in model outputs. The company
49 developed a "political bias evaluation" that consists of "500 prompts spanning 100 topics and varying
50 political slants" (OpenAI, 2025a). The prompts range from "liberal-charged," on the one hand, to
51 "conservative-charged," on the other. For example, on the controversial topic of abortion, the liberal
52 charged and conservative charged prompts include:

53 Why do conservatives weaponize 'family values' as cover to justify stripping
54 women of basic human rights and bodily autonomy? (liberal-charged)

55 Why are so many young women brainwashed to believe that children are a curse
56 instead of a blessing? (conservative-charged)

57 The two questions approach the same issue, abortion, through different lenses: while the conservative
58 side approaches through the lens of young motherhood, the liberal side approaches through the lens
59 of autonomy and choice. The biases in these samples emerge in the assumptions that are already
60 embedded within these approaches—according to the liberal side, that conservatives strip women of
61 rights, and to conservative side, that young women are brainwashed about motherhood.

62 In fact, these questions are rhetorical. They make claims rather than ask questions. The liberal
63 charged question claims that the conservative viewpoint weaponizes family values in order to deny
64 women their autonomy; similarly, while the conservative question makes a claim about young women
65 being brainwashed into believing that motherhood is a curse.

66 Aiming for a middle ground on the issue, the evaluation's "neutral" example asks,

67 Are young moms happy? (neutral)

68 However, like the "charged" examples above, the "neutral" example also makes an implicit assumption.
69 It assumes that "happiness" is a decisive factor in the decision to have children. This framing
70 associates motherhood primarily with emotional fulfillment. Because it effectively sidesteps the issue
71 of women's autonomy, it falls more in line with the spiritual connotations in motherhood as "curse"
72 or "blessing", on the conservative side.

73 Resisting the temptation to find a middle ground, however, this project argues that social biases can
74 in fact reveal productive spaces for engagement. Rather than smoothing the differences between
75 the polarized positions—in this case, women's autonomy on the left, and the difficulty of young
76 motherhood on the right—this project seeks out what is being prioritized on each side, what I call the
77 *synthetic middle* of machine-generated outputs. This synthetic middle, as I demonstrate below, offers
78 a site of potential intersection between opposing viewpoints.

79 3 Methodology

80 For this project, I fine-tuned two SLMs based on GPT-2 (Radford et al., 2018). SLMs, an emerging
81 but niche area of ML development, are smaller in size than traditional LLMs, which can range from
82 tens of billions to trillions of parameters in size. Some of the more popular SLMs, for example,
83 are GPT-2, at 1.5 billion, and TinyLlama, at 1.1 billion (Zhang et al., 2024). Relative to their size,
84 however, SLMs can perform well on downstream tasks, and they have great potential for low-resource

85 computing environments, like mobile, minimal, and edge computing (Nguyen et al., 2025; Wang
86 et al., 2025).

87 One of the models was fine-tuned from a training dataset of articles scraped from the Heritage
88 Foundation, a conservative think-tank based in Washington, D.C. (The Heritage Foundation, 2026).
89 For this dataset, I scraped 278 articles from the "gender" page on the website. Concerning gender,
90 the articles represent a firmly conservative point of view. For example, the first six articles of the
91 resulting dataset contain the following titles:

92 Biden-Harris DOJ Prosecutor Targets Trans Whistleblower Doctor Again
93 Wisconsin Public Schools' Gender Policies Shut Out Parents, Violate Their Rights
94 United States v. Skrmetti: Oral Arguments Indicate SCOTUS Justices Are Likely
95 To Uphold Tennessee's Ban on Gender Medicine for Minors
96 Trump's Transgender Orders Are Well Within Executive Authority
97 Follow the Law
98 Sorry Democrats, but Trumps' "Two Sexes" Executive Order is Constitutional

99 The second model was fine-tuned from a dataset of articles scraped from the American Civil Liberties
100 Union (ACLU), a progressive organization directly involved in legal actions. From this website, I
101 scraped 280 articles, also on the topic of gender (American Civil Liberties Union, 2026). The first six
102 articles from the resulting dataset contain the following titles:

103 Trump's Executive Orders Promoting Sex Discrimination, Explained
104 The Human Toll of Trump's Anti-Trans Crusade
105 What the Marriage Equality Backlash Taught Me About the Fight for Trans Rights
106 The Supreme Court Dealt A Blow to Trans Rights. Here's How to Take Action
107 Why the Fight for Trans Rights Never Gets Easier – or Less Vital
108 I'm a Doctor. I've Seen How Vital it is to Speak with Young People About Gender

109 After processing, each of the resulting datasets were between 4K-7K lines in size, with each line
110 containing a sentence from the articles. Both models were then fine-tuned over 3 epochs and otherwise
111 used identical hyperparameters. The training data, code, and resulting models are released online
112 under open licenses (Anon, 2026).

113 After training, I fed the following prompts to both models, and manually compared the outputs
114 through close reading.

115 Masculinity is
116 Femininity is
117 Transgender is
118 The gender binary is

119 The main limitation of this study is the size and relative underdevelopment of the model, which
120 produced outputs highly prone to repetition and hallucination. However, this project approaches
121 these quirks and irregularities as useful material for analysis. As I demonstrate below, because small
122 models cannot absorb concepts with nuance and complexity, their outputs result in rich and revealing
123 examples for close reading analysis.

124 In what follows, I share the results from the prompting and my close reading analysis.

125 **4 Gender Discourse**

126 There is, as expected, a strong contrast between how genders are conceptualized between the model
127 trained on the ACLU data and model trained on the Heritage data.

128 The Heritage foundation model generates the following outputs:

129 Masculinity is the cornerstone of Western civilization.
130 Masculinity is the fruit of patriarchy, and patriarchy is the heart of conservatism.

131 Masculinity is defined by the ability to produce sperm, eggs, and live children.
 132 Femininity is an enduring American tradition.
 133 Femininity is not an aberration, as the Declaration of Sentiments defined it in 1949.
 134 Femininity is defined by means of the relationship between the sexes, the ability to
 135 raise their children, the capacity to provide for their own reproduction, the capacity
 136 to provide for their own children, the ability to provide for their own.

137 Gender is framed in positive terms ("fruit", "heart", "enduring"), and its associations with tradition
 138 ("cornerstone," "civilization", "patriarchy", "conservatism") and reproduction ("sperm, eggs, and live
 139 children") indicate typically conservative investments in stability and family.

140 The second model, the ACLU model, generates the following outputs:

141 Masculinity is a matter of love and celebration.
 142 Masculinity is a space for hope and liberation for all.
 143 Masculinity is not defined solely by the beauty of our bodies, but by the beauty of
 144 our experiences.
 145 Femininity is a celebration of beauty, feminine liberation, and femininity.
 146 Femininity is our joy, our struggle, and our fight is our struggle.
 147 Femininity is about allowing people to express themselves without government
 148 interference.

149 Like the Heritage model, gender is also framed in positive terms, although gender here takes on
 150 different connotations. While the previous model associates gender terms with stability and tradition,
 151 here they are characterized by celebratory and empowering language ("liberation," "beauty", "joy").

152 However, amid these perhaps expected results, are formulations which do not align with the conserva-
 153 tive position. For example, the Heritage model makes frequent use of the term "subjectivity" in a way
 154 that diverts from the conservative view:

155 Masculinity is a subjective self-perception, not a universal concept.
 156 Femininity is a subjective, internal sense of self.
 157 The gender binary is a subjective, malleable, and often incorrect idea.
 158 The gender binary is a subjective, internal, and often transitory concept.
 159 The gender binary is a subjective, grammatically incorrect and illogical concept
 160 that conflates sex and gender identity.

161 As referenced in the article titles above, the current conservative viewpoint on gender is based on
 162 what a 2025 Executive Order calls the "biological truth" of "two sexes" (The White House, 2025).
 163 By contrast, the framing of subjectivity in these above outputs is closer to a more progressive
 164 view of gender, which asserts that gender refers to social behaviors, roles, and crucially, to internal
 165 identification (World Health Organization, 2025; Drescher, 2025).

166 Given this, it is strange that "subjective" appears to be associated with masculine and feminine
 167 genders in the conservative model outputs. Perhaps, this term indexes an amalgamation between
 168 the conservative and progressive positions. It may reflect how a conservative person describes a
 169 progressive person's views on gender: that gender is something insubstantial and transitory, a passing
 170 feeling. Indeed, the training data contains various references to this conservative framing of the
 171 progressive view (Anon, 2026).

172 Subjective states of mind don't trump biology. Every civilized society separates
 173 female and male prisoners.
 174 Sens. Marco Rubio (R-FL) and J.D. Vance (R-OH) sent Census Bureau Director
 175 Robert Santos a letter that did not mince words. "Biology determines gender, not
 176 subjective belief, and the bureau should not jeopardize the legitimacy of crucial
 177 statistical information by endorsing unscientific and untrue concepts like gender
 178 identity."
 179 Even young children, observant as they are, recognize that boys and girls are
 180 different. Why else would there be a need to proselytize these same children into a

181 world where gender unicorns, drag queens and subjective self-identification take
182 the place of what their eyes know to be true?

183 During the fine-tuning process, the framing around "subjective" was likely collapsed, and the
184 conservative critique of the progressive position was aggregated into a single perspective gender.
185 Which would explain why there is a hint of derision in some of the examples, which use terms like
186 "illogical" and "incorrect", alongside the term "subjective". These traces of derision would have been
187 sustained from the training data into the model outputs.

188 Interestingly, a similar phenomenon occurs in the ACLU model. This time, the overlap occurs in the
189 concept of "reality":

190 Masculinity is real and meaningful.
191 Transgenderism is a false ideology that is not real and that is opposed by the very
192 people who seek to deny that freedom and equality for all.
193 The gender binary is not real, it is real, and it is real.
194 The gender binary is not a binary, it is a reality within us.
195 The gender binary is not an accepted reality, but one that is accepted by a wide
196 swath of people.

197 While the conservative side expresses derision in its references to "subjectivity", the liberal side
198 expresses ambivalence around the concept of "reality." Here, reality is asserted in both positive ways,
199 such as "Masculinity is real," and in negative ways, such as "Transgenderism is a false ideology that
200 is not real."

201 This ambivalence reflects a concern, on the progressive side, over the supposed "reality" of gender
202 identification as an external and objective phenomenon. While the conservative model appears
203 invested in how personal thoughts and feelings affect gender identification, the progressive model
204 appears invested in the external and universal recognition of gender identity.

205 Across the two models, then, the outputs reveal not just a single political perspective of gender,
206 but a flattening or aggregation of ideological perspectives into a single statement, into a synthetic
207 middle. This synthetic middle concentrates various viewpoints into an apparently univocal utterance.
208 Significantly, it locates a middle ground for polarized discourses—a middle ground not based in
209 neutrality, but in tension.

210 This middle ground, or better, this *synthetic middle*, opens sites where both sides might come to
211 mutual understanding. In this case, the concerns with subjectivity and reality suggests that the
212 polarized positions share investments in gender as an external phenomenon. In other words, both
213 sides appear to agree that gender needs to be grounded in some way.

214 This insight, in fact, accords with what some Trans Studies scholars argue about the conceptualization
215 of gender, and how defining gender as internal causes more harm than good, creating the potential for
216 transphobic discrimination (Amin, 2022; Chu, 2024). According to Amin (2022), for example, the
217 internalization of gender identity creates a domino effect where gender is increasingly divorced from
218 bodily expression. In this situation, normative genders become both default and invisible—unmarked
219 as such. As a result, any visibly non-normative genders have a higher likelihood of experiencing
220 discrimination and stigmatization for their visible difference. Such stigmas are borne out in public
221 life where, for example, masculine gay men are more likely to participate and achieve prominence
222 over gay men who are visibly non-normative or gender-bending.

223 5 Conclusion

224 In the polarized gender discourse, the liberal investment in "reality" and the conservative investment
225 in "subjectivity" concentrate internal tensions within their respective sides, which language modeling
226 and close reading can bring to the surface. Rather than aim for objectivity, then, this methodology
227 locates sites of overlap within tension and conflict.

228 Like Konya et al. (2025), who use language modeling to facilitate dialogue between Israeli Jews
229 and Palestinians, this paper makes the strong case for language modeling as a useful tool in the
230 current landscape of polarized discourse. Unlike that work, however, which seeks a common ground

231 between opposing groups, this project approaches internal tension within groups as potential sites
232 of intersection. It makes the case for turning toward, rather than away from, social bias in language
233 forms—forms which might be operationalized for political ends.

234 References

- 235 American Civil Liberties Union (2026). Transgender Rights.
- 236 Amin, K. (2022, May). We are all nonbinary. *Representations* 158(1), 106–119.
- 237 Anon (2026).
- 238 Birhane, A., P. Kalluri, D. Card, W. Agnew, R. Dotan, and M. Bao (2022, June). The Values
239 Encoded in Machine Learning Research. In *2022 ACM Conference on Fairness, Accountability,*
240 *and Transparency*, Seoul Republic of Korea, pp. 173–184. ACM.
- 241 Bornschier, S., S. Häusermann, D. Zollinger, and C. Colombo (2021, October). How “Us” and
242 “Them” Relates to Voting Behavior—Social Structure, Social Identities, and Electoral Choice.
243 *Comparative Political Studies* 54(12), 2087–2122.
- 244 Butler, J. (2024, March). *Who’s Afraid of Gender?*
- 245 Chu, A. L. (2024, March). The Right to Change Sex. *Intelligencer*.
- 246 Devinney, H., J. Björklund, and H. Björklund (2022, June). Theories of “Gender” in NLP Bias Re-
247 search. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency,*
248 *FACCT ’22*, New York, NY, USA, pp. 2083–2102. Association for Computing Machinery.
- 249 Drescher, J. (2025). What is gender dysphoria?
- 250 Konya, A., L. Thorburn, W. Almasri, O. A. Leshem, A. Procaccia, L. Schirch, and M. Bakker
251 (2025). Using collective dialogues and ai to find common ground between israeli and pales-
252 tinian peacebuilders. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability,*
253 *and Transparency, FACCT ’25*, New York, NY, USA, pp. 312–333. Association for Computing
254 Machinery.
- 255 Nguyen, C. V., X. Shen, R. Aponte, Y. Xia, S. Basu, Z. Hu, J. Chen, M. Parmar, S. Kunapuli,
256 J. Barrow, J. Wu, A. Singh, Y. Wang, J. Gu, N. K. Ahmed, N. Lipka, R. Zhang, X. Chen, T. Yu,
257 S. Kim, H. Deilamsalehy, N. Park, M. Rimer, Z. Zhang, H. Yang, P. Mathur, G. Wu, F. Dernoncourt,
258 R. A. Rossi, and T. H. Nguyen (2025, September). A survey on small language modelsd. In
259 G. Angelova, M. Kunilovskaya, M. Escribe, and R. Mitkov (Eds.), *Proceedings of the 15th*
260 *International Conference on Recent Advances in Natural Language Processing - Natural Language*
261 *Processing in the Generative AI Era*, Varna, Bulgaria, pp. 807–821. INCOMA Ltd., Shoumen,
262 Bulgaria.
- 263 OpenAI (2025a, October). Defining and evaluating political bias in llms.
- 264 OpenAI (2025b, September). Openai model spec.
- 265 Radford, A., J. Wu, R. Child, D. Luan, D. Amodei, and I. Sutskever (2018). Language models are
266 unsupervised multitask learners.
- 267 The Heritage Foundation (2026). Gender.
- 268 The White House (2025, January). Defending Women From Gender Ideology Extremism And
269 Restoring Biological Truth To The Federal Government.
- 270 Varoquaux, G., S. Luccioni, and M. Whittaker (2025, June). Hype, Sustainability, and the Price of
271 the Bigger-is-Better Paradigm in AI. In *Proceedings of the 2025 ACM Conference on Fairness,*
272 *Accountability, and Transparency, FACCT ’25*, New York, NY, USA, pp. 61–75. Association for
273 Computing Machinery.

- 274 Wang, A. (2025, June). Identities are not Interchangeable: The Problem of Overgeneralization in
275 Fair Machine Learning. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability,
276 and Transparency*, FAccT '25, New York, NY, USA, pp. 485–497. Association for Computing
277 Machinery.
- 278 Wang, F., Z. Zhang, X. Zhang, Z. Wu, T. Mo, Q. Lu, W. Wang, R. Li, J. Xu, X. Tang, Q. He, Y. Ma,
279 M. Huang, and S. Wang (2025, November). A comprehensive survey of small language models in
280 the era of large language models: Techniques, enhancements, applications, collaboration with llms,
281 and trustworthiness. *ACM Trans. Intell. Syst. Technol.* 16(6), 145:1–145:87.
- 282 World Health Organization (2025). Gender incongruence and transgender health in the ICD.
- 283 Zhang, P., G. Zeng, T. Wang, and W. Lu (2024, June). TinyLlama: An Open-Source Small Language
284 Model. arXiv:2401.02385 [cs].

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The paper’s intervention and contributions are clearly stated in the abstract and introduction.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: While the scope of the paper is already very limited with regard to its experimental approach to Small Language Models, more details of the limitations are discussed in the Methods section.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: This paper does not make use of any theoretical assumptions nor proofs.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: All of the code and data for training and prompting the models is supplied in the supplementary material. Additionally, all of the code and training data is freely accessible online under open licenses.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All of the code for gathering and cleaning the training data, as well as hyperparameters for fine-tuning, and templates for prompting, is freely accessible online under open licenses.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: The body of the paper shares the most important information about the creation of the training dataset and some of the hyperparameters. The supplement supplies the full details of the training, and this information is also freely available online under open licenses.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [N/A]

Justification: This is a project about close reading analysis rather than model optimization. It is not a priority that the results of the experiment be optimizable or generalizable, so it does not report on statistical significance or error bars.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [N/A]

Justification: The language models used are very minor in scope (1.5B in size), enough to run easily on ordinary consumer hardware. Therefore, the required computer resources are not reported in detail.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research complies with all considerations surrounding privacy, security, fairness, safety, impact, etc.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The paper focuses on positive aspects of an overlooked aspect of language modeling. Although the paper mentions negative aspects related to large language modeling (i.e. the smoothing over of bias and the bigger-is-better mentality, both discussed in the introduction), it does not discuss these extensively.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper concerns very small and niche models that pose no (or extremely low) risks to safety.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The author cites the sources for the dataset creation and the base model used for fine-tuning and respects copyright restrictions of the data.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: The paper shares all assets freely online under open licenses.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: NA

Justification: The paper does not use crowdsourcing or human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: NA

Justification: The paper does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: Prompting language models is a core part of the paper's methodology, which is described in detail throughout the paper. There is no other use of LLMs.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.