

Unmarking Gender: Trans Studies Interventions in NLP

Abstract

This paper brings theorizing from Trans Studies to bear on recent NLP methods for detecting and evaluating anti-trans bias in Large Language Models. Drawing in particular on Trans Studies' critique of gender identification, and its emphasis on gender as a visible, expressive, and socially material phenomenon, the paper examines how treating gender primarily as an internal identity label limits what bias detection systems are able to perceive. Systems that conceptualize gender as internal render dominant genders unmarked, while obscuring the embodied behaviors and presentations that are most frequently targeted by discrimination. Through close analysis of two recent frameworks for studying trans bias—the TIDEs and EQUITABL approaches (Lameiro et al. 2025, Devinney et al. 2024)—the paper shows how different bias detection and mitigation methods might draw from the insights of Trans Studies theorizing. The paper concludes by outlining directions for future research that integrate theoretical insights from Trans Studies to better capture the material and relational dimensions of gender bias in language technologies.

Introduction

There is a global urgency to define gender as a category. In various countries around the world, gender designations like "transgender" and "nonbinary" are prompting backlash that has resulted in the retraction of hard-gained rights for LGBTQ+ groups, such as the ban on gender-affirming care for minors in the USA (Karni and Kavi, 2025). This backlash takes what it terms "gender ideology", or nontraditional views of gender that are not based on biological sex, as a rallying cry for political

priorities that are similarly based in xenophobic logics, such as anti-immigrant discrimination (Human Rights Research Center, 2025).

In response to attacks on what it calls "gender ideology," some progressive groups have doubled down on a definition of gender that is based on internal identification rather than biological reality (Dresher 2025). Drawing from seminal theories of gender, such as Judith Butler's theory of Gender Performativity (1990, discussed in more detail below), this perspective of gender breaks down the rigid and causal relationship between sex and gender, leaving room for individuals to identify in ways that may or may not correspond to their assigned sex at birth.

The same assumptions that drive this global controversy over defining gender also operate in methods for studying transgender bias in language, such as in the field of NLP bias detection and mitigation techniques. Through the analysis of concrete linguistic features, NLP methods attempt to make explicit the ways that language forms encode perceptions on gender. Over the past several years, such methods have ranged from Word Embedding approaches (Bolukbasi et al. 2016), to Coreference Resolution tasks (Rudinger et al. 2018), and more recently, Language Generation methods (Kaneko et al. 2024).

Across the disciplinary divide, in the meantime, the field of Trans Studies has been exploring paradigms that can adequately describe and represent all forms of gender manifestations. Within these conversations, one perspective, called "trans materialism", is based on the notion that gender ought to be conceived as a social and behavioral phenomenon (Amin 2022). Its logic states that if gender is conceived as an outward phenomenon, instead of an internal identity, then it will be easier to make visible those who are non-conforming, and therefore most marginalized, by gender norms.

This paper explores how these conceptualizations from trans materialism might be applied to current methods in NLP, focusing on two recent methods for detecting transgender bias in language, TIDEs and EQUITABL (Lameiro et al. 2025, Devinney et al. 2024). In a close and deep reading of their approaches and insights, this analysis draws these methods' potential sites of engagement with

Trans Studies theory. It argues that despite their contrasting approaches, with one focusing on binary classification and the other on qualitative analysis, both methods significantly gain from a consideration of Trans Studies theorizing on gender bias. Specifically, the emphasis on marking visible and relational aspects of gender has direct application to labelling and qualitative coding methods in NLP.

A central part of this paper's motivation is its "close reading" analytical method, which it demonstrates in practice. It argues that a deep and close analysis of methodology, called close reading in Humanities disciplines, can draw out significant alignments between diverse methodologies and fields of study. This paper therefore contributes to work that calls for and brings interdisciplinary theorizing into NLP, such as from the context of this conference (Wang et al. 2025, Birhane et al. 2022).

Gender ideology

The term "gender ideology", which the Southern Poverty Law Center classifies as a derogatory term, refers to the view that LGBTQ+ identity is an internal and false belief system (SPLC 2023). In various countries around the world, gender ideology has risen into a national issue, serving as a rallying cry for various political groups to garner support. In the European continent, far-right groups have used this concept to prop up other forms of social discrimination, such as those against migrants. In Italy, for example, gender fearmongering has been linked to a larger white nationalist campaign (Indelicato and Lopes 2024). Alternatively, on the African continent, similar arguments have been used for contrasting ends, in anti-colonial discourses, in which gender ideology is framed as a Western import that threatens traditional gender roles. In Uganda, in particular, the past few years have seen a wave of anti-LGBT+ legislation, with one law penalizing "aggravated homosexuality" with the death penalty (BBC 2023).

In the USA, concerns over gender ideology have prompted a dramatic increase in legislation which limits the rights of trans people. Over the last five years, the number of proposed anti-trans bills each year has surpassed the last, with 2025 proposing over 1000 such bills in legislative houses across the country (Trans Legislation Tracker 2025). Additionally, from the White House, a mounting wave of transphobia has propelled a series of Executive Orders with titles that frame trans rights as a defense of womens' rights: "Defending Women From Gender Ideology Extremism And Restoring Biological Truth To The Federal Government" and "Keeping Men Out of Women's Sports". These initiatives double down on the primacy of sex as "binary and biological" (The White House 2025a, The White House 2025b). Like the anti-immigrant tactics that turn white supremacist and nationalist logics of discrimination, these anti-trans initiatives draw from the patriarchal beliefs that so-called "biological" women need explicit protection from transwomen.

Proponents of gender ideology-based arguments aim their critiques on the progressive view of gender. In the abovementioned Executive Orders, for example, gender ideology is described as "an ever-shifting concept of self-assessed gender identity, permitting the false claim that males can identify as and thus become women and vice versa, and requiring all institutions of society to regard this false claim as true." (The White House 2025a).

Indeed, the progressive view tends to define gender as a subjective or personal perception. The American Psychiatric Association, for example, defines gender identity as "a person's inner sense of being a girl/woman, boy/man, some combination of both, or something else" ("What is Gender Dysphoria?" 2025). Similarly, the World Health Organization defines gender identity as "a person's innate, deeply felt internal and individual experience of gender," and contrasts it to biological sex, adding that gender identity "may or may not correspond to the person's physiology or designated sex at birth" ("What is Gender Dysphoria?" 2025, "Gender and health" 2025).

Some Tran Studies scholars posit the progressive definition of gender is the problem. That the push to disentangle gender from sex and to solidify sex as "binary and biological" is a reaction to the

ambiguity of how gender is conceptualized in trans affirming discourse. Andrea Long Chu (2024), for example, points this out that, "The left must reckon with its part in this. It has hung trans rights on the thin peg of gender identity".

The view of gender as an internal phenomenon thus inspires transphobic reasoning around the issue of medical care, which argues that a physical transition is not necessary. Chu, a trans woman, explains that the pro-trans movement has "largely failed to form a coherent moral account of *why* someone's gender identity should justify the actual biological interventions that make up gender-affirming care." In other words, Chu asks, if gender really is something that is internal, then why should a gender non-conforming or transgender person need to altar any external characteristics?

Gender categories

Contemporary, progressive notions of gender descend from a widely accepted view that gender is a social phenomenon rather than a biological reality. This view can be traced to Judith Butler's seminal theory of gender performativity (1991), which dissolves the traditionally strict and causal relation between sex, gender, and sexuality. Butler's theory stipulates that gender is not based on a biological reality, nor does it determine sexuality (with male gendered people preferring female gendered people). Rather, Butler argues that gender is an effect produced social norms and expectations, and that it manifests as a series of behaviors. According to Butler's theory, then, gender becomes real not in any essential or ontological sense, but through its active and external expression (1991).

In the 35 years since Butler's theory, the progressive definition of gender has evolved into the notion of identity. Moving from a kind of behavior (one that, Butler is careful to point out, is not at will, but deeply compelled by social rules), gender takes on an internal aspect of identification. Gender as identity, an internal feeling or identification as male, female, a combination of the two, or

nonbinary genders, thus diverges from gender as expression, which denotes the masculine, feminine, or androgynous traits that are visibly apparent and enacted. Furthermore, in this conception, sex is still understood to be a bodily component, though one not based firmly in a binary, nor does it necessarily prioritize any biological marker, like the presence of external genitals, sex hormones or chromosomes.

Trans Studies scholar Kadjia Amin (2022) refers to this splitting of internal identity from external appearance or expression as a "divergence model". The split in this understanding of gender has, according to Amin, created a domino effect where gender is increasingly idealized and divorced from bodily expression.

For example, Amin explains, the divergence model drives the split between the categories of heterosexual and homosexual, in which the homosexual category originally encompassed all kinds of sex, gender, and sexuality non-conformity. This view accords with sexual historiographers who, tracing the categories further back into the 19th century, explain that heterosexuality as a category is a response to the creation of the category of homosexual [Foucault 1978, Halperin 2002]: in this view, heterosexuality arose as a reaction to work by 19th century sexologists, such as that by Richard von Krafft-Ebing (1893) and Havelock Ellis (1900), who defined homosexuality primarily as a psychological pathology.

Just as heterosexuality emerged as a response to homosexuality, Amin (2022) argues, so did the concept of cisgender, which responds to the emergence of the term "transgender" in the 1990s. Whereas the concept of homosexuality used to incorporate gender-variant forms of sexuality, for example, forms of male femininity such as the "queen", and of female masculinity in the form of the "butch", the emergence of transgender offered a new category that was then used to hold "all manner of gender-bending" (Amin 2022). One particular effect of the emergence of transgender is to create an idealized notion of normativity that does not accurately describe any gender or sexuality in lived reality.

As a result of these splits, Amin argues, terms like heterosexual, cisgender, and binary take up "increasingly idealized and uninhabitable normative categories" (Amin 2022). Gender identity becomes buried "deeper and deeper within the private recesses of the self, where it increasingly disavows any relation to the social" (Amin 2022). It becomes transported into an internal and idealized form, and retains the privilege of being unnoticed as such. This idealization, according to Amin, has enabled the LGBTQ+ movement to campaign for rights and recognition over the past several decades.

The effect of such idealized conceptualizations of gender and/or sexuality is to re-stigmatize non-normative gender expressions. Designating non-normative gender as an *identity* that is invisible, that is not immediately apparent despite being marked, effectively marginalizes gender when it appears as *expression*. Amin gives a common example from homosexual culture, where gay men who are normatively masculine, "masc" or "butch" men, have higher social status than feminized gay men, "queens" and "fairies", due to the cultural capital of masculinity (2022).

One way to resist the effects of the divergence model, stigmatizes non-normative genders that are externally expressed, is to embrace gender as external. Which is to say that defining gender as an external expression or behavior is more likely to make visible and validate all genders.

Trans bias in NLP

I now turn to my close reading of two bias detection and mitigation techniques. For each of these techniques, I examine how their approaches attend to gender as a relational quality.

Framework #1: TIDEs

The first Trans Bias detection method is a project called **Transphobia Identification in Digital Environments**, or TIDEs (Lameiro et al. 2025). TIDEs comprises a dataset and model created to

detect anti-trans bias in language. Building from the framework of a "trans technology," a "technology that addresses the unique needs and challenges faced by trans people and communities" (Haimson 2025), this tool centers trans expertise in the detection and mitigation of transphobic language.

To that end, TIDEs has two explicit goals: first, to create a dataset representing varied forms of transphobia, in particular forms that are more difficult to detect, like "TERF". Referring to "Trans Exclusionary Radical Feminism," TERF perspectives are difficult to detect because they identify as feminist despite negating the rights of transwomen and girls. The second goal is to center trans and nonbinary experience and expertise in the data labelling work, in what the authors call "care-centered data practices" (Lameiro et al. 2025). The project chose labellers who personally identify as trans or nonbinary, and who "have personal experience with encountering and detecting online transphobia", to label social media posts as "transphobic or "not transphobic".

The dataset comprises social media posts, which are scraped from Tumblr, Truth Social, YouTube, and Reddit, labelled with "transphobic" and "not transphobic" labels. For example, the post, "Transgender females are actually biological males. They should not be allowed to compete in womens' sports. Period!" is labelled as transphobic, while the post, "Transgender Swimmer Lia Thomas Booted From Olympics After Losing Legal Case", is labelled as not transphobic.

As the researchers explain, labelling transphobia brings its own challenges, when many of its manifestations appear through forms of irony and other rhetorical contexts than be difficult to detect. For example, the following text was labelled as "not transphobic", despite the majority of the language actually represents a transphobic point of view:

Hey guys, I just wanted to get your opinion on this:

Whenever I meet someone I really quickly try to establish whether or not they have a dick, I don't just restrict this to potential sexual partners I mean

everyone I meet. It just seems super important you know? Like how am I supposed to sleep at night unless I catalogue everyone's genitalia?

[...] I've started flat out asking everyone I meet. I try to be polite, like I look at the ground and draw it out all "Sooo do you like.... Have a dick?" but here is the thing: people look at me like I'm crazy for asking! What the hell is up with that? I even had one girl tell me I was a totally creep!

Does this stuff happen to you?

/ok transjerkers I'm done for a little while, fuck it's nice to vent

Because this text, on the surface, parrots a transphobic perspective, it could be rightly labelled as such. Detecting this post correctly as "not transphobic" relies on recognizing its irony which, although present in the tone of enforced sincerity throughout the post, is only confirmed at the end, with the assertion that "it's nice to vent". Another clue is the reference to "transjerkers", an in-group moniker for a trans-dominated subreddit that emphasizes parody, "r/transcirclejerk". Therefore, the only hints that this text is *not* transphobic are so subtle that they could be easily missed, particularly by those who lack the in-group cultural knowledge and language cues.

Contrast this example with a similar one, which uses a sincere tone to make a serious request. This post, also drawn from reddit, asks for advice on gender presentation:

Hey! I'm going through some gender experimentation, so I'd like to tap into masculinity for awhile(usually [sic] I identify as just pure androgynous, but I struggle with pulling that off). So, to begin with, I'll try to master visual/behavioral androgyny first. I have long hair but don't want to cut it(a lot of manly men have long hair anyway), so how could I fix it up to look more masculine? Also, how could I dress more masculine-like? I don't have a very

manly looking body shape, but I'd like to dress to give the illusion of one.

Basically, any and all advice on how to appear androgynous/masculine is appreciated.

This post, which is labelled as "not transphobic," shares with the previous "transphobic" example a tone that appears sincere and supplicating. The difference has to do with an external rhetorical situation that is not clearly visible in the actual post. In other words, context matters whether something is transphobic.

The close and careful data labelling work, which was aided by a detailed coding schema, distinguishes this strategy from others. Although it sorts examples into a binary classification, the coding guidelines in fact enables labellers to identify and mark various kinds of transphobia. This approach demonstrates how detecting transphobia requires attention not only to surface-level language but also to cultural and rhetorical cues. By centering trans and nonbinary expertise through care-centered data practices, the project addresses limitations of conventional content moderation systems that may otherwise miss transphobia's subtle cues.

Framework #2: EQUITABL

EQUITABL, the **Explore, Query, & Understand Implicit Textual Bias in Language data** project (Devinney et al. 2024) uses a mixed-methods approach to analyze social bias in language, which deliberately addresses bias toward trans and nonbinary identities. This project combines text generation, topic modelling, and close reading methodologies to examine "stereotypes, hegemonic norms, [and] erasure of identity" across identity groups (Devinney et al. 2024).

Unlike TIDEs, which is a binary classification method, this method emphasizes a qualitative approach. Built from a method the researchers developed in a previous paper which proposes a "combination of quantitative and qualitative methods, where the quantitative part produces a view of the data of a size suitable for qualitative analysis" (Devinney et al. 2020). For the this more recent

version of the project, the researchers expanded their methodology to include text generation, with the main analysis focusing on close reading a set of LLM-generated narratives.

In brief, the process involves prompting both English and Swedish language models using a set of templates designed to generate narratives about various identity groups. The resulting narratives are then fed to a topic model, which will organize them into topics. Finally, the researchers extract the top narratives for each topic and proceed with close reading analysis, which is the main focus of the experiment.

The templates for prompting the language models consist of three components: first, an imperative phrase asking the model to "write a story about" or "write a description of", followed by the subject of the narrative, and closing with a predicate (Devinney et al. 2024). The templates cycle through different imperatives, racial and gender identities, religions, and predicates to create a varied set of prompts. Possible variations include, "a Black Christian at the doctor's office," "a trans Muslim with a special skill," "a white nonbinary person with tears running down their face." The result is to create a diverse set of stories with protagonists spanning various subject positions, represented in both positive, negative, and neutral contexts.

The study's central intervention is its close reading methodology, which it includes in the place of quantitative scoring. As the researchers explain, they intentionally use qualitative metrics in order to "encourage users to critically engage with the topics and patterns... rather than merely being incentivized to chase 'better' scores" (Devinney et al. 2024). This metric consists of a series of questions that the researcher uses to guide close reading, which include, "How are people and bodies described?" and "Which stereotypes are used?", as well as questions that relate gender to specific themes.

According to the researchers, close readings indicate that narratives centering feminine identities are mostly linked to unpleasant emotions like depression and anxiety, a quality that is exacerbated in narratives about trans women. Narratives about masculine subjects, by contrast, are mostly

positive in tone, associated with happy events like weddings and parties. Additionally, narratives about nonbinary persons express overt themes of struggle and the need for community support, and the researchers note that these struggles tend to end on a positive, "uplifting" note.

Amid these perhaps unsurprising results, researchers note an unexpected pattern of refusal from the model, which declined to discuss certain topics and identity groups. In these refusals, the models explain that their programming instructions require them to either stay away from topics of identity that may promote "negative or derogatory language" (Devinney et al. 2024). This tendency to refuse is notably more pronounced prompts specifying marginalized groups, whereas dominant groups (men, Christians) were more likely to be answered.

However, there was an exception to the pattern of refusals as they concerned socially dominant groups. When identity markers from socially dominant groups were combined in an additive way, such as "white", "Christian", and "man" to form "a white Christian man", the model was more likely to refuse than similarly intersectional, though less dominant, identity combinations. For example, a model would be more likely refuse a prompt about a "white Christian man" than a "Black christian man." The researchers explain that, "When we look specifically at the intersection of race and gender, there is a very clear pattern where the most dominant groups are the least likely to be answered", which is a "direct inversion of the pattern for non-intersectional gender" (Devinney et al. 2024). For example, they explain, the least refused prompt was the phrase "a man," while the most refused one was "a white Christian" (Devinney et al. 2024).

Of course, refusing to answer a prompt about a specific identity group is its own form of discrimination. As the researchers rightly point out, such refusals implicitly reinforce that only a select and dominant subset of identities are fit for discussion, while other, more marginalized groups are not: "The refusals themselves actually produce risk and harm. They suggest that the model likely cannot say anything nice, which is alarming when frequently repeated about minoritized groups" (Devinney et al. 2024).

This notable inversion of the usual pattern also speaks to another peculiarity of the results, that which the researchers call the "unmarked majority." A lack of identity specification in the model prompts will lead the model to assume a dominant identity, which is a white, straight, cisgender, Christian male identity. The researchers explain that, "When not otherwise specified, 'a person' is assumed to be a man, as well as likely white, straight, cisgender, and Christian; additionally [...] being a provider" (Devinney et al. 2024).

While perhaps unsurprising, this aspect of the results may perhaps explain why the model refuses to generate narratives of intersectional dominant groups ("a white Christian"). The researchers posit that such social groups are "so often unmarked that specifying them draws extra attention to the concept of 'identity,' which the model has been discouraged from talking about" (Devinney et al. 2024).

Discussion

In what follows, I discuss how these two gender bias evaluation frameworks engage with Trans Studies' theorizing about external gender expression and how it relates to stigma. As Amin (2022) argues, due to the historical tendency of LGBTQ+-related categories to distinguish internal and idealized identities from visibly different ones, a kind of stigmatization accompanies non-normative genders that are externally expressed. Given this focus, how does gender as external presentation or behavior operate in the above frameworks for studying gender bias in NLP?

The close examination of these two frameworks for detecting trans bias—TIDEs and EQUITABL—demonstrate how bias evaluation can reveal concrete aspects of gender bias in language. Despite their contrasting approaches, TIDEs relying on binary classification and EQUITABL on qualitative analysis, both frameworks reveal the importance of context and interpretation in identifying transphobia. TIDEs highlights how subtle rhetorical cues, specifically

irony, and in-group knowledge complicate automated detection, while EQUITABL shows how generative models reproduce hegemonic norms through refusal and the dominant, unmarked subject.

It is important to remember that existing evaluations of LLMs show that the the most frequent and default subject category is the straight white male, and that other subject categories, like woman, black, and homosexual, are less prevalent in both training data and model outputs (Guo and Caliskan 2021). Additionally, research has also shown a correlation between the presence and frequency of certain identities in the training data and their association with positive and negative terms (Devinney 2024, van Loon 2022): van Loon et al. (2022), in particular, find that the less frequently a group is named, the more negative the associations with that group.

If a group's relatively frequent presence the training data has a positive effect on their representation, one might reasonably assume that making marginalized subjects *more frequent* in the training data, to be on par with the frequency of representation of dominant subjects, may mitigate bias about these groups in model outputs. However, this assumption doesn't take into account that the default subject is often unmarked. As Devinney et al. (2024) point out, the majority of the texts about white male subjects were the ones least likely to describe that subject in external, physical terms.

That is perhaps because of an inverse relationship between physical markers of identity and positions of social dominance, which Devinney et al.'s (2024) work bear out. Here, less physical description is itself a marker of social dominance. In those results, white women are described in physical terms only slightly more than white men, while marginalized identities are described in more detail than both groups. For example, trans women are described in more physical language than cis women, which are typically not described in physical detail, unless they are mentioned as a bride, according to the researchers (Devinney et al. 2024). These results suggest that the more that a person is described in physical terms, the more likely they come from a marginalized group.

One solution, then, may be to increase physical descriptions of all subjects—though especially white, male, heterosexual, Christian identified subjects. Some work already exists toward this end, by Reif and Schwartz (2023). In their article, "Fighting Bias with Bias," the researchers share a promising approach: *amplifying* rather than reducing bias in a model's training dataset. As they point out, "filtering [out bias] can obscure the true capabilities of models to overcome biases, which might never be removed in full from the dataset" [Reif and Schwartz 2023]. Instead of an approach that reduces bias, they offer one of increasing it, such as by including phrases like "the pretty doctor" so that a model will interpret "doctor" to be female.

Toward this end, this paper makes one overall suggestion: that researchers spend more time with close reading methods, such as those displayed by TIDEs and EQUITABL methods. Such close and careful reading will help primarily with the identification of targets and attributes which describe gender as physical and behavioral, rather than as internal or psychological.

This suggestion applies to multiple approaches to gender bias evaluation. The TIDEs researchers, for example, developed a robust qualitative coding system that could catch various examples of transphobia. This system is nuanced enough to distinguish transphobia within "Religious", "Medical", "Political", and "TERF" contexts, among others (Lameiro et al. 2025). Although the final model is a binary classification model, these labels are built into the detection process through the data labellers' careful coding work.

Amin's point—that if gender were conceptualized as an outward, expressive phenomenon, rather than an identity category, there would be less social stigma—thus applies to widely different methods for evaluating gender and specifically trans bias in NLP. Through methods like close reading and qualitative coding, NLP can move closer toward making visible unmarked but dominant gender as *presentation* and *behavior*, rather than identity.

Conclusion

This paper has shown that, while contemporary gender frameworks that conceptualize gender as an internal identity, detached from outward expression and embodiment, has enabled important forms of recognition and rights-based advocacy, it also produces unintended consequences: idealized and invisible norms (such as cisgender and heterosexuality) are reinforced, while visibly non-normative gender expressions become further stigmatized. The problem with the universal, unmarked subject is precisely that: despite not being marked, he is overrepresented. This overrepresentation does not result from the persistent or explicit labels as such. The white man does not carry a descriptive label as such, he simply is.

Incorporating this critique from Trans Studies offers a productive direction for future work in the study of gender bias in NLP. Rather than treating gender solely as an abstract identity variable, NLP research benefits directly from interdisciplinary work that attend to embodiment, behavior, and visibility, making dominant norms explicit rather than naturalized.

Future work in NLP ought to draw more deliberately on bias research that is grounded in theoretical texts from other fields. As demonstrated by Devinney et al. (2024), engaging explicitly with frameworks of gender (Butler 1990), intersectionality (Crenshaw 1991), and the matrix of domination (Collins 2002) enables a more nuanced understanding of how implicit dominance structures emerge in language forms. Similarly, Lameiro et al. (2025) show how drawing from an ethics of care (de la Bellacasa 2011) can inform methodological choices, particularly in data labelling practices. Together, these approaches illustrate how theoretically grounded interdisciplinary work can deepen NLP research on bias.

Most importantly, however, this paper has shown how close reading methods—that is, sustained, close attention to language—can help materialize how gender operates in NLP systems. This paper's

central argument can be boiled down to the logic that if gender is a expressive and relational phenomenon, and language has the power to shape social structures, then reading gender bias ought to also be a material practice.

Generative AI Disclosure Statement

Generative AI was not used for this paper.